

# Docker File , Networking & Swarm

## 1. Docker file

Dockerfile defines what goes on in the environment inside your container. Access to resources like networking interfaces and disk drives is virtualized inside this environment, which is isolated from the rest of your system, so you need to map ports to the outside world, and be specific about what files you want to “copy in” to that environment

(both Dockerfile and dockerfile would be accepted, but not DockerFile, DOckerfile)

### Docker file Component:

#### **FROM**

The base image for building a new image. This command must be on top of the dockerfile.

#### **MAINTAINER**

Optional, contains the name of the maintainer of the image.

#### **RUN**

Used to execute a command during the build process of the docker image.

#### **ADD**

Copy a file from the host machine to the new docker image. There is an option to use an URL for the file, docker will then download that file to the destination directory.

```
1  # fetch node v4 LTS codename argon
2  FROM node:argon
3
4  # Request samplename build argument
5  ARG samplename
6
7  # Create app directory
8  RUN mkdir -p /usr/src/spfx-samples
9  WORKDIR /usr/src/spfx-samples
10
11 #Install app dependencies
12 RUN git clone https://github.com/SharePoint/sp-dev-fx-webparts.git .
13 WORKDIR /usr/src/spfx-samples/samples/$samplename
14
15 # install gulp on a global scope
16 RUN npm install gulp -g
17
18 # RUN ["npm", "install", "gulp"]
19 RUN npm install
20 RUN npm cache clean
21
22 # Expose required ports
23 EXPOSE 4321 35729 5432
24
25 # Run sample
26 CMD ["gulp", "serve"]
27
```

**ENV**

Define an environment variable.

**CMD**

Used for executing commands when we build a new container from the docker image.

**ENTRYPOINT**

Define the default command that will be executed when the container is running.

**WORKDIR**

This is directive for CMD command to be executed.

**USER**

Set the user or UID for the container created with the image.

**VOLUME**

Enable access/linked directory between the container and the host machine.

## **Best practices for writing Dockerfiles:**

Docker can build images automatically by reading the instructions from a Dockerfile, a text file that contains all the commands, in order, needed to build a given image. Dockerfiles adhere to a specific format and use a specific set of instructions

- **Containers should be ephemeral:** mean that it can be stopped and destroyed and a new one built and put in place with an absolute minimum of set-up and configuration
- **Use a .dockerignore file:** To exclude files which are not relevant to the build, without restructuring your source repository, use a .dockerignore file.
- **Avoid installing unnecessary packages:** To exclude files which are not relevant to the build
- **Each container should have only one concern:** Decoupling applications into multiple containers makes it much easier to scale horizontally and reuse containers.
- **Minimize the number of layers**

### **Some important links:**

[https://docs.docker.com/v17.09/engine/userguide/eng-image/dockerfile\\_best-practices/#build-cache](https://docs.docker.com/v17.09/engine/userguide/eng-image/dockerfile_best-practices/#build-cache)

<https://docs.docker.com/v17.09/engine/reference/builder/>

## Push images to Docker Hub:

```
docker login --username=<yourhubusername>

docker tag <imageid> yourhubusername/<reponame>:<tagname>

docker push yourhubusername/<reponame>

docker pull yourhubusername/<reponame>:<tagname>
```

Docker Hub Account: push docker images to the hub

## Docker File Practice:

Create an empty directory **app**. Change directories (cd) into the new directory, create a file called Dockerfile, copy-and-paste the following content into that file, and save it.

```
# Use an official Python runtime as a parent image
FROM python:2.7-slim

# Set the working directory to /app
WORKDIR /app

# Copy the current directory contents into the container at /app
ADD . /app

# Install any needed packages specified in requirements.txt
RUN pip install --trusted-host pypi.python.org -r requirements.txt

# Make port 80 available to the world outside this container
EXPOSE 80

# Define environment variable
ENV NAME World

# Run app.py when the container launches
CMD ["python", "app.py"]
```

Create two more files called requirements.txt and app.py (When the above Docker file is built into an image, app.py and requirements.txt is present because of that Docker file's ADD command, and the output from app.py is accessible over HTTP thanks to the EXPOSE command)

## requirements.txt

Flask

Redis

## app.py

```
from flask import Flask
from redis import Redis, RedisError
import os
import socket

# Connect to Redis
redis = Redis(host="redis", db=0, socket_connect_timeout=2, socket_timeout=2)

app = Flask(__name__)

@app.route("/")
def hello():
    try:
        visits = redis.incr("counter")
    except RedisError:
        visits = "<i>cannot connect to Redis, counter disabled</i>"

    html = "<h3>Hello {name}!</h3>" \
        "<b>Hostname:</b> {hostname}<br/>" \
        "<b>Visits:</b> {visits}"
    return html.format(name=os.getenv("NAME", "world"), hostname=socket.gethostname(), visits=visits)

if __name__ == "__main__":
    app.run(host='0.0.0.0', port=80)
```

## How to build docker images from Dockerfile

```
docker build -t myfirstdocker .  
docker images  
docker run -p 4001:80 myfirstdocker or docker run -itd -p 4000:80 myfirstdocker  
docker ps  
docker ps -a  
docker start 03e9a3566b60
```

---

## 2. Docker Networking:

When we install docker by default it will create 3 networks.

docker network ls

```
[root@ip-172-31-30-165 ec2-user]# docker network ls  
NETWORK ID          NAME                DRIVER              SCOPE  
b469e26c310e        bridge             bridge              local  
f940a7666910        host               host                local  
8bb20e494c5b        none               null                local
```

**Bridge:** Bridge plays very important role for whole networking systems in docker container.

**Host:** Host has very important role for only for docker host machine. It doesn't have responsibility to do networking with containers.

**Null:** Null network can be used when we don't want to allow any container to use network. Container cannot do any communication with any devices.

**Default Network:** The Bridge network represents the docker0 network present in all Docker installations. To change the network run below command

```
docker run --network=<NETWORK>
```

Syntax: docker run -it -d --network=none --name node6 ubuntu:latest /bin/bash

```
[root@ip-172-31-30-165 ec2-user]# docker run -it -d --network=none --name node6
ubuntu:latest /bin/bash
293d188457175a7426ec1c38e8e83ac7c8859d8a29c7066d4cf16cda89f99b81
[root@ip-172-31-30-165 ec2-user]# docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED
293d18845717	ubuntu:latest	"/bin/bash"	22 seconds ago

```
Up 22 seconds
node6
```

When you go inside your container and try to install anything, it will not work because here we have attached null network to our container.

```
[root@ip-172-31-30-165 ec2-user]# docker exec -it 293d18845717 /bin/bash
root@293d18845717:/# apt-get update
Err:1 http://security.ubuntu.com/ubuntu xenial-security InRelease
Temporary failure resolving 'security.ubuntu.com'
Err:2 http://archive.ubuntu.com/ubuntu xenial InRelease
Temporary failure resolving 'archive.ubuntu.com'
Err:3 http://archive.ubuntu.com/ubuntu xenial-updates InRelease
Temporary failure resolving 'archive.ubuntu.com'
Err:4 http://archive.ubuntu.com/ubuntu xenial-backports InRelease
Temporary failure resolving 'archive.ubuntu.com'
Reading package lists... Done
W: Failed to fetch http://archive.ubuntu.com/ubuntu/dists/xenial/InRelease Temporary failure resolving 'archive.ubuntu.com'
W: Failed to fetch http://archive.ubuntu.com/ubuntu/dists/xenial-updates/InRelease Temporary failure resolving 'archive.ubuntu.com'
W: Failed to fetch http://archive.ubuntu.com/ubuntu/dists/xenial-backports/InRelease Temporary failure resolving 'archive.ubuntu.com'
W: Failed to fetch http://security.ubuntu.com/ubuntu/dists/xenial-security/InRelease Temporary failure resolving 'security.ubuntu.com'
W: Some index files failed to download. They have been ignored, or old ones used instead.
root@293d18845717:/#
```

## Network Linking containers:

```
docker pull redis
docker images
docker run -d --name redis1 redis
```

```
[root@ip-172-31-30-165 ec2-user]# docker pull redis
Using default tag: latest
latest: Pulling from library/redis
b0568b191983: Pull complete
6637dc5b29fe: Pull complete
7b4314315f15: Pull complete
67b22db27e51: Pull complete
350dbcc91819: Pull complete
eee5ee716895: Pull complete
Digest: sha256:26c93c5b06aaa323bb1089500f42b0dd158138772348b865e364127f1d554982
Status: Downloaded newer image for redis:latest
[root@ip-172-31-30-165 ec2-user]# docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
redis	latest	4760dc956b2d	2 days ago	107MB
pranik/modifyimage	latest	bf41b3e14a13	2 days ago	177MB
myimage	latest	4d16e26172d3	2 days ago	112MB
pranik/myimageonhub	latest	4d16e26172d3	2 days ago	112MB
ubuntu	latest	f975c5035748	9 days ago	112MB
tomcat	latest	108db0e7c85e	4 weeks ago	558MB
httpd	latest	01154c38b473	4 weeks ago	177MB
hello-world	latest	f2a91732366c	3 months ago	1.85kB

```
[root@ip-172-31-30-165 ec2-user]# docker run -d --name redis1 redis
6e10e073c2b7822a08bab3de53ba6fc593d5a6adc0955f4938e765e46e5ce469
[root@ip-172-31-30-165 ec2-user]# docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
6e10e073c2b7	redis	"docker-entrypoint.s..."	6 seconds ago	Up 5 seconds	6379/tcp	redis1

```
docker run -it --link redis1:redis --name redisclient1 busybox
```

communication happens with a secure channel

```
/ # cat /etc/hosts
127.0.0.1      localhost
::1           localhost ip6-localhost ip6-loopback
fe00::0       ip6-localnet
ff00::0       ip6-mcastprefix
ff02::1       ip6-allnodes
ff02::2       ip6-allrouters
172.17.0.3     redis 6e10e073c2b7 redis1
172.17.0.4     b6e518d2b2fc

/ # ping redis
PING redis (172.17.0.3): 56 data bytes
64 bytes from 172.17.0.3: seq=0 ttl=255 time=0.079 ms
64 bytes from 172.17.0.3: seq=1 ttl=255 time=0.068 ms
64 bytes from 172.17.0.3: seq=2 ttl=255 time=0.064 ms
64 bytes from 172.17.0.3: seq=3 ttl=255 time=0.066 ms
64 bytes from 172.17.0.3: seq=4 ttl=255 time=0.067 ms
64 bytes from 172.17.0.3: seq=5 ttl=255 time=0.065 ms
64 bytes from 172.17.0.3: seq=6 ttl=255 time=0.065 ms
^C
--- redis ping statistics ---
7 packets transmitted, 7 packets received, 0% packet loss
round-trip min/avg/max = 0.064/0.067/0.079 ms
/ #
```

Redis client(busybox) can easily communicate with the redis server

## Create Custom (User Defined) Bridge Network:

1. User-defined bridges provide better isolation and interoperability between containerized applications.
2. User-defined bridges provide automatic DNS resolution between containers:
  - Containers on the default bridge network can only access each other by IP addresses, unless you use the `--link_option`, which is considered legacy. On a user-defined bridge network, containers can resolve each other by name or alias.

```
docker network create my-bridge-net  
docker network ls
```

### we also can define subnet ranges for user defined bridge network

```
docker network create --driver=bridge --subnet=192.168.0.0/16 my-bridge-net2
```

### Delete custom default network

```
Docker network rm my-bridge-net2
```

### If you want to run new container to custom bridge network

```
docker create --name my-nginx --network my-bridge-net --publish 6080:80 nginx:latest
```

Login to your existing ubuntu container and Install ping

```
apt-get install iputils-ping
```



Now try to ping any container which is on different network and same bridge network (EX: nginx container which we created in previous step and old ubuntu containers. And verify ping)

**Note: Only Bridge can be created.**

We CAN'T create host and none network like bridge network

```
docker network create --driver=host myhost
```

(If we try to create then Docker will throw error:

Error response from daemon: only one instance of "host" network is allowed)

---

## Docker log directory:

```
/var/lib/docker
```

## Docker troubleshooting:

```
docker inspect <container Id>
```

```
docker inspect <image Id>
```

```
docker logs <container id>
```

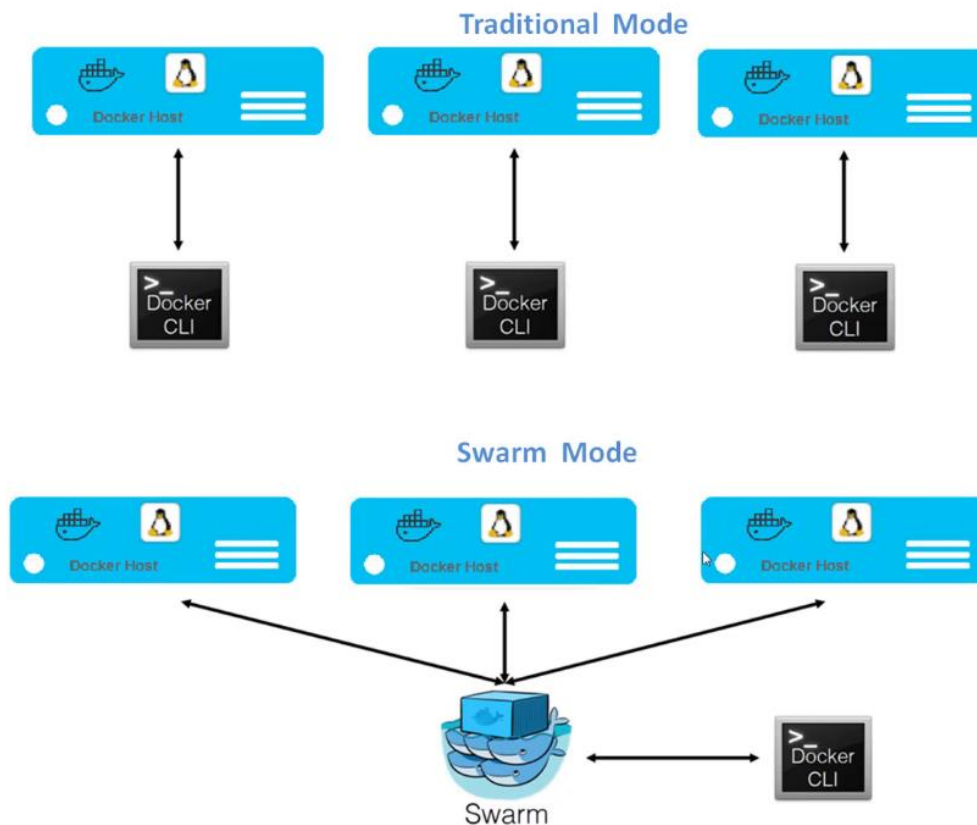
```
docker logs --since 2018-05-25 <container id>
```

### 3. Docker Swarm

Swarm is native clustering for the Docker. In the context of swarm, a cluster is a pool of Docker hosts that acts as a bit like a single large docker host. You can also run swarm services and standalone containers on the same Docker instances.

The cluster management and orchestration features embedded in the Docker Engine are built using swarmkit. Swarmkit is a separate project which implements Docker's orchestration layer and is used directly within Docker.

A swarm consists of multiple Docker hosts which run in **swarm mode** and act as managers (to manage membership and delegation) and workers (which run swarm services). A given Docker host can be a manager, a worker, or perform both roles



## **Features of Swarm:**

- **Cluster management integrated with Docker Engine:** Use the Docker Engine CLI to create a swarm of Docker Engines where you can deploy application services. You don't need additional orchestration software to create or manage a swarm.
- **Decentralized design:** Instead of handling differentiation between node roles at deployment time, the Docker Engine handles any specialization at runtime. You can deploy both kinds of nodes, managers and workers, using the Docker Engine. This means you can build an entire swarm from a single disk image.
- **Declarative service model:** Docker Engine uses a declarative approach to let you define the desired state of the various services in your application stack. For example, you might describe an application comprised of a web front end service with message queuing services and a database backend.
- **Scaling:** For each service, you can declare the number of tasks you want to run. When you scale up or down, the swarm manager automatically adapts by adding or removing tasks to maintain the desired state.
- **Desired state reconciliation:** The swarm manager node constantly monitors the cluster state and reconciles any differences between the actual state and your expressed desired state. For example, if you set up a service to run 10 replicas of a container, and a worker machine hosting two of those replicas crashes, the manager creates two new replicas to replace the replicas that crashed. The swarm manager assigns the new replicas to workers that are running and available.
- **Multi-host networking:** You can specify an overlay network for your services. The swarm manager automatically assigns addresses to the containers on the overlay network when it initializes or updates the application.
- **Service discovery:** Swarm manager nodes assign each service in the swarm a unique DNS name and load balances running containers. You can query every container running in the swarm through a DNS server embedded in the swarm.
- **Load balancing:** You can expose the ports for services to an external load balancer. Internally, the swarm lets you specify how to distribute service containers between nodes.

- **Secure by default:** Each node in the swarm enforces TLS mutual authentication and encryption to secure communications between itself and all other nodes. You have the option to use self-signed root certificates or certificates from a custom root CA.
- **Rolling updates:** At rollout time you can apply service updates to nodes incrementally. The swarm manager lets you control the delay between service deployments to different sets of nodes. If anything goes wrong, you can roll-back a task to a previous version of the service.

### **On Manager Node**

1. Execute the following command with the manager nodes IP for initializing the swarm cluster.

```
docker swarm init
```

OR

```
docker swarm init --advertise-addr MANAGER
```

Run output command to your swarm nodes (to add all the nodes to cluster mode)

To know the swarm cluster info

```
docker info
```

To know the information about all the nodes in the cluster,

```
docker node ls
```

---IMP commands

docker swarm leave --- leave the swarm

docker swarm update --update the swarm

Promote a node (worker become manager) - `docker node promote docker03`

Demote a node (manager become worker) - `docker node demote docker03`

Docker swarm join-token worker --- to get worker token